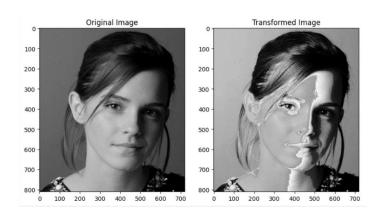
In20-S5-EN3160 - Assignment 1

Index Number – 200709K

Wickramanayake R.S.D

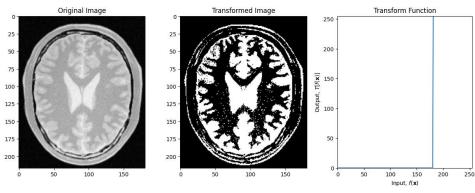
Github Repository: - RuchchaSD/EN3160 (github.com)

```
%matplotlib inline
import cv2 as cv
import numpy as np
import matplotlib.pyplot as plt
# Prepare Transform Function
t1 = np.linspace(0, 50, 51).astype('uint8')
t2 = np.linspace( 100, 255, 150 - 50).astype('uint8')
t3 = np.linspace(150,255 ,255-150).astype('uint8')
transform = np.concatenate((t1, t2), axis=0).astype('uint8')
transform = np.concatenate((transform, t3), axis=0).astype('uint8')
#open original image
img_orig = cv.imread('emma.jpg', cv.IMREAD_GRAYSCALE)
#transform image
image_transformed = cv.LUT(img_orig, transform)
#plot results
fig, ax = plt.subplots(1,3, figsize=(15, 5))
ax[0].imshow(img orig, cmap='gray')
ax[0].set_title('Original Image')
ax[1].plot(transform)
ax[1].set_xlabel(r'Input, $f(\mathbf{x})$')
ax[1].set_ylabel(r'Output, $T[f(\mathbf{x})]$')
ax[1].set_xlim(0, 255)
ax[1].set_ylim(0, 255)
ax[2].imshow(image_transformed, cmap='gray')
ax[2].set_title('Transformed Image')
plt.show()
```



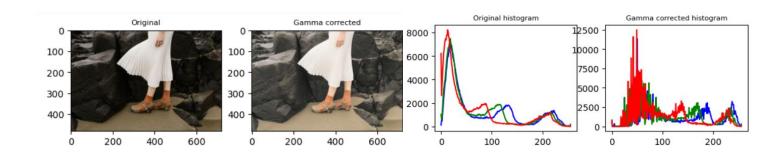
Question 2 (Only Image transformation code snippet)

```
t1 = np.linspace(0, 0, 181).astype('uint8')
t2 = np.linspace( 255, 255, 255 - 180).astype('uint8')
transform = np.concatenate((t1, t2), axis=0).astype('uint8')
#open original image
img_orig = cv.imread('BrainProtonDensitySlice9.png', cv.IMREAD_GRAYSCALE)
#transform image
image_transformed = cv.LUT(img_orig, transform)
```



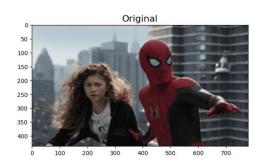
Question 3 (Only Image transformation code snippet)

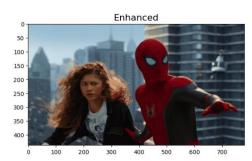
```
# read the image
img_orig = cv.imread('highlights_and_shadows.jpg', cv.IMREAD_COLOR)
# convert to LAB color space
img_gamma = cv.cvtColor(img_orig, cv.COLOR_BGR2LAB)
img_L = img_gamma[:,:,0]
# apply gamma correction to the L channel
gamma = 2 # A gamma value of 2 gave a visually pleasing output
table = np.array([(i/255.0)**(1.0/gamma)*255.0 for i in
np.arange(0,256)]).astype('uint8')
img_L_gamma = cv.LUT(img_L, table)
img_gamma[:,:,0] = img_L_gamma
img_gamma = cv.cvtColor(img_gamma, cv.COLOR_LAB2RGB)
img_orig = cv.cvtColor(img_orig, cv.COLOR_BGR2RGB)
```

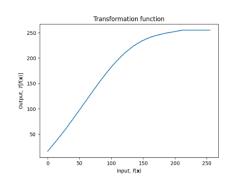


Question 4 (Only Image transformation code snippet)

```
img_orig = cv.imread('spider.png', cv.IMREAD_COLOR)
# Convert img_orig to HSV
img_hsv = cv.cvtColor(img_orig, cv.COLOR_BGR2HSV)
# Get the saturation channel
img s = img hsv[:,:,1]
# Define the transformation function
def transform(x, a, sigma):
    return min(x + a * 128 * np.exp(-((x - 128) ** 2) / (2 * sigma ** 2)), 255)
a = 0.7
sigma = 70
# Generate the LUT using the transformation function
table = np.array([transform(i, a, sigma) for i in np.arange(0,
256)]).astype('uint8')
# Apply the LUT to the saturation channel
img s transformed = cv.LUT(img s, table)
#recombine image
img_hsv[:,:,1] = img_s_transformed
img_enhanced = cv.cvtColor(img_hsv, cv.COLOR_HSV2BGR)
# Plot original and transformed image
f, axarr = plt.subplots(1, 2, figsize=(15, 10))
axarr[0].imshow(cv.cvtColor(img_orig, cv.COLOR_BGR2RGB))
axarr[0].set_title('Original', size=16)
axarr[1].imshow(cv.cvtColor(img_enhanced, cv.COLOR_BGR2RGB))
axarr[1].set_title('Enhanced', size=16)
plt.show()
fig, ax = plt.subplots()
ax.plot(table)
ax.set_xlabel(r'Input, $f(\mathbf{x})$')
ax.set_ylabel(r'Output, $T[f(\mathbf{x})]$')
ax.set title('Transformation function')
```





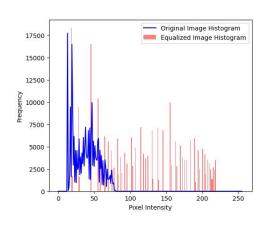


Question 5 (Only Image transformation code snippet)

```
input_image = cv.imread("shells.tif", cv.IMREAD_GRAYSCALE)
# Calculate histogram
rows, cols = input image.shape
histogram = np.zeros((256,), dtype=np.uint16)
for i in range(rows):
    for j in range(cols):
        intensity = input_image[i, j]
        histogram[intensity] += 1
# calculate cdf
cdf = np.zeros((256,), dtype=np.uint16)
for i in range(256):
    for j in range(i + 1):
        cdf[i] += histogram[j] * (255 / (rows * cols))
    cdf[i] = round(cdf[i], 0)
cdf = cdf.astype(np.uint16)
output_image = np.zeros_like(input_image)
for i in range(rows):
    for j in range(cols):
        intensity = input image[i, j]
        output_image[i, j] = cdf[intensity]
```



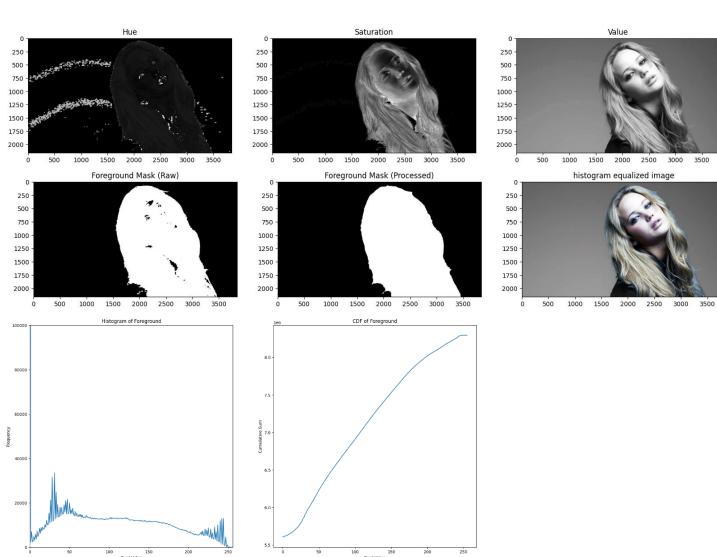




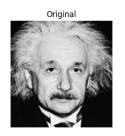
Question 6 (Only Image transformation code snippet)

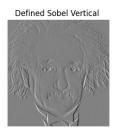
```
img = cv2.imread('jeniffer.jpg', cv2.IMREAD_COLOR)
img_hsv = cv2.cvtColor(img, cv2.COLOR_BGR2HSV)
h, s, v = cv2.split(img_hsv)
# Define a threshold value for saturation
threshold_value = 20
# Create a binary mask for the foreground based on the saturation channel
foreground_mask = np.zeros_like(s)
foreground_mask[s > threshold_value] = 255
foreground_raw = foreground_mask.copy()
foreground_mask = cv2.morphologyEx(foreground_mask, cv2.MORPH_CLOSE,
cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (80, 80)))
foreground_mask = cv2.erode(foreground_mask,
cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (20, 20)), iterations=1)
background_mask = cv2.bitwise_not(foreground_mask)
```

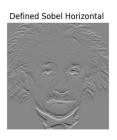
```
# Apply the mask to image to extract foreground objects
img= cv2.cvtColor(img, cv2.COLOR BGR2RGB);
foreground r = cv2.bitwise and(img[:,:,0], foreground mask)
foreground_g = cv2.bitwise_and(img[:,:,1], foreground_mask)
foreground b = cv2.bitwise_and(img[:,:,2], foreground mask)
img_foreground = cv2.merge((foreground_b, foreground_g, foreground_r))
img_foreground = cv2.cvtColor(img_foreground, cv2.COLOR_BGR2GRAY)
#get the histogram of the foreground
hist_foreground = cv2.calcHist([img_foreground],[0],None,[256],[0,256])
#calculate the cumulative distribution function
cdf_foreground = hist_foreground.cumsum()
#histogram equalize the foreground
foreground r = cv2.equalizeHist(foreground r)
foreground_g = cv2.equalizeHist(foreground_g)
foreground_b = cv2.equalizeHist(foreground_b)
img_masked = cv2.merge((foreground_r, foreground_g, foreground_b))
background_mask = cv2.bitwise_not(foreground_mask)
#extract the background
background_r = cv2.bitwise_and(img[:,:,0],background_mask)
background_g = cv2.bitwise_and(img[:,:,1],background_mask)
background_b = cv2.bitwise_and(img[:,:,2],background_mask)
background = cv2.merge((background_b, background_g, background_r))
img new = cv2.add(img masked, background)
```



```
# write a function to fiter the given image with given kernel
def filterImg(image, kernel):
   img_h, img_w = image.shape
   kernel_h, kernel_w = kernel.shape
   pad h = kernel h // 2
   pad w = kernel w // 2
   image float = cv.normalize(image.astype('float'), None, 0.0, 1.0,
cv.NORM MINMAX)
   output = np.zeros(image.shape, 'float')
   for y in range(pad h,img h-pad h):
       for x in range(pad w,img w-pad w):
            output[y, x] = np.dot(image_float[y - pad_h:y + pad_h + 1, x -
pad_w:x + pad_w + 1].flatten(), kernel.flatten())
    return output
# Read image
img = cv.imread('einstein.png', cv.IMREAD GRAYSCALE)
# Sobel horizontal
kernel_h = np.array([(-1, -2, -1), (0, 0, 0), (1, 2, 1)], dtype='float')
# Sobel vertical
kernel_v = np.array([(-1, 0, 1), (-2, 0, 2), (-1, 0, 1)], dtype='float')
# Using Propertry of Convolution
kernel_1 = np.array([(1), (2), (1)], dtype='float')
kernel_2 = np.array([(1, 0, -1)], dtype='float')
# Using OpenCV
imgcv = cv.filter2D(img,-1,kernel_v)
imgch = cv.filter2D(img,-1,kernel_h)
# Using Defined Function
imgv = filterImg(img, kernel_v)
imgh = filterImg(img, kernel_h)
# Using Property of Convolution
imgpv = cv.filter2D(img, -1, kernel_1)
imgpv = cv.filter2D(imgpv, -1, kernel 2)
```

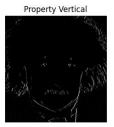








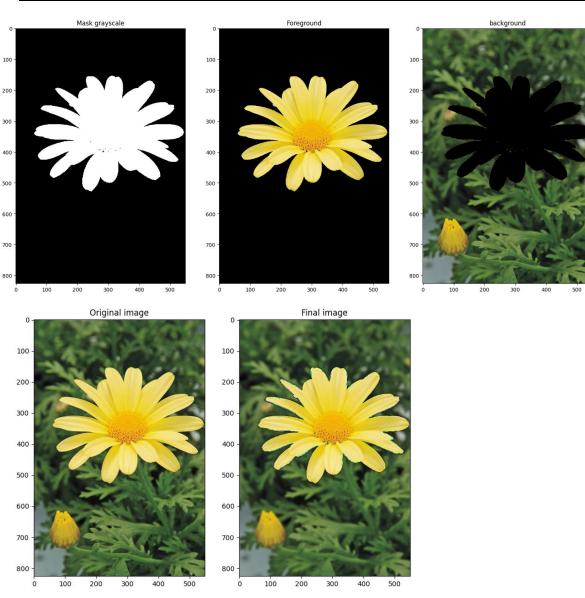




```
def zoom nearest neighbor(image, zoom factor):
    height, width, _ = image.shape
    new_height = int(height * zoom_factor)
    new_width = int(width * zoom_factor)
    new_image = np.zeros((new_height, new_width, 3), dtype=np.uint8)
    for y in range(new_height):
        for x in range(new width):
           src_x = int(x / zoom_factor)
           src_y = int(y / zoom_factor)
           new_image[y, x] = image[src_y, src_x]
    return new image
def zoom bilinear(image, zoom factor):
    height, width, _ = image.shape
    new_height = int(height * zoom_factor)
    new_width = int(width * zoom_factor)
    new_image = np.zeros((new_height, new_width, 3), dtype=np.uint8)
    for y in range(new_height):
        for x in range(new_width):
           src_x = x / zoom_factor
           src_y = y / zoom_factor
            x1, y1 = int(src_x), int(src_y)
            x2, y2 = x1 + 1, y1 + 1
            if x2 >= width:
               x2 = width - 1
            if y2 >= height:
               y2 = height - 1
            dx = src_x - x1
            dy = src_y - y1
            new_image[y, x] = (
                (1 - dx) * (1 - dy) * image[y1, x1] +
                dx * (1 - dy) * image[y1, x2] +
                (1 - dx) * dy * image[y2, x1] +
                dx * dy * image[y2, x2]
            ).astype(np.uint8)
    return new_image
```

```
Image[ 1 ] ssd_nearest : 31.284316486625514 ssd_bilinear : 39.257033179012346
Image[ 2 ] ssd_nearest : 11.902013310185184 ssd_bilinear : 16.21177662037037
Image[ 3 ] ssd_nearest : 17.171342909907853 ssd_bilinear : 22.376757387099232
Image[ 4 ] ssd_nearest : 78.73781724215534 ssd_bilinear : 81.65874063625257
Image[ 5 ] ssd_nearest : 50.57724609375 ssd_bilinear : 53.70764858217593
Image[ 6 ] ssd_nearest : 30.553995949074075 ssd_bilinear : 35.51769129372428
Image[ 7 ] ssd_nearest : 27.964142659505207 ssd_bilinear : 30.22005997721354
Image[ 8 ] ssd_nearest : 14.37148410910911 ssd_bilinear : 19.219769269269268
Image[ 9 ] ssd_nearest : 21.148100694444445 ssd_bilinear : 26.6669085590277778
Image[ 10 ] ssd_nearest : 21.42844935276991 ssd_bilinear : 25.21551477169893
Image[ 11 ] ssd_nearest : 94.7471351111111 ssd_bilinear : 94.56599555555556
```

```
rect = (50,150,500,380)
bgdModel = np.zeros((1,65),np.float64)
fgdModel = np.zeros((1,65),np.float64)
mask = np.zeros(img.shape[:2],np.uint8)
cv.grabCut(img,mask,rect,bgdModel,fgdModel,5,cv.GC_INIT_WITH_RECT)
mask2 = np.where((mask==2)|(mask==0),0,1).astype('uint8')
foreground = img * mask2[:, :, np.newaxis]
background = img - foreground
blurred_background = cv.GaussianBlur(background, (13, 13), 0)
final_image = blurred_background + foreground
final_image = cv.cvtColor(final_image, cv.COLOR_BGR2RGB)
```



Q9.(c) The dark edges were introduced due to gaussian blur applied to the background image. The part were the foreground flower was present was black when gaussian blur was applied. So the kernel for gaussian blur would have taken the black pixels in the cut part for calculating values for the edges in the final image.